

# Searching and Analyzing Data with Elasticsearch: Getting Started

---

## INTRODUCING ELASTICSEARCH



**Janani Ravi**

CO-FOUNDER, LOONYCORN

[www.loonycorn.com](http://www.loonycorn.com)

# Overview

**A little search engine history and the importance of search**

**Basics steps involved in indexing and searching documents**

**The inverted index, the heart of a search engine**

**An introduction to Elasticsearch and its basic building blocks**

**Set up and install Elasticsearch on your local machine and check cluster health**

# What You Need in Your Toolkit

---



# Prerequisites

**Familiarity with the command line on a Mac, Linux or Windows machine**

**Familiarity with using RESTful APIs to perform actions**

**A very basic understanding of distributed computing**



# Install and Setup

The latest version of Elasticsearch, 5.4.0 requires Java version 8

A Mac, Linux or Windows machine on which Elasticsearch can be installed



# Course Overview

**Introduction** to basic concepts in Elasticsearch, download and install

**Building** an index, **adding** documents to it both individually and in bulk

**Search** queries on an index using the Query DSL

**Analysis** of data on an index using aggregations

# A Brief History of Search

---

# Brief History of Search

**1945**

Vannevar Bush first talks of the need to index records

**1991**

Tim Berners-Lee combined hypertext, TCP and DNS to imagine WWW

**1993**

Excite improved search by using statistical analysis of word relationships

**1970s**

The ARPANet network which laid the foundation of the modern internet

**1993**

Primitive search engines, linear search of URLs, very basic ranking

**1994**

Yahoo offered a directory of useful webpages i.e. a portal



# Brief History of Search

**1994**

Lycos provided ranking relevance, prefix matching, a huge catalog

**1996**

Inktomi pioneered the paid inclusion model

**1998**

Google ranking pages based on how many other pages link to it

**1994**

Altavista had natural language queries, inbound link checking

**1997**

ask.com had natural language search, human editors for queries


**Today**

Google, Bing, Baidu, Naver, Yahoo

# How Does Search Work?

---

# What Is the Objective of Search?



Find the **most relevant** documents  
with your search terms

# Most Relevant Document for Search Terms



**Know of the  
document's  
existence**



**Index the  
document for  
lookup**



**Know how  
relevant the  
document is**



**Retrieve  
ranked by  
relevance**

# Most Relevant Document for Search Terms



**Web crawler**



Index the  
document for  
lookup



Know how  
relevant the  
document is



Retrieve  
ranked by  
relevance

# Most Relevant Document for Search Terms



Web crawler



**Inverted  
index**



Know how  
relevant the  
document is



Retrieve  
ranked by  
relevance

# Most Relevant Document for Search Terms



Web crawler



Inverted  
index



**Scoring**



Retrieve  
ranked by  
relevance

# Most Relevant Document for Search Terms



Web crawler



Inverted  
index



Scoring



**Search**



# Most Relevant Document for Search Terms



**Web crawler**



**Inverted  
index**



**Scoring**



**Search**

# Search Is Not Restricted to The Web

---

# Sites Have Their Own Search



**E-commerce**



**Video**



**E-learning**

# The Inverted Index

---

# Documents Have Content

**House Stark**

**Winter is coming**

**House Baratheon**

**Ours is the fury**

**House Tyrell**

**Growing Strong**

# Tokenize Text into Words

winter
is
coming
ours
the
fury
growing
strong

**split words**

**lowercased**

**removed  
punctuation**

# Tokenize Text into Words

<b>winter</b>		<b>1</b>	
<b>is</b>		<b>2</b>	
<b>coming</b>		<b>1</b>	
<b>ours</b>		<b>1</b>	
<b>the</b>		<b>1</b>	
<b>fury</b>		<b>1</b>	
<b>growing</b>		<b>1</b>	
<b>strong</b>		<b>1</b>	

# Tokenize Text into Words

<b>winter</b>	<b>1</b>
<b>is</b>	<b>2</b>
<b>coming</b>	<b>1</b>
<b>ours</b>	<b>1</b>
<b>the</b>	<b>1</b>
<b>fury</b>	<b>1</b>
<b>growing</b>	<b>1</b>
<b>strong</b>	<b>1</b>



# Tokenize Text into Words

winter	1	Stark
is	2	Stark, Baratheon
coming	1	Stark
ours	1	Baratheon
the	1	Baratheon
fury	1	Baratheon
growing	1	Tyrell
strong	1	Tyrell

# Tokenize Text into Words

winter	1	Stark
is	2	Stark, Baratheon
coming	1	Stark
ours	1	Baratheon
the	1	Baratheon
fury	1	Baratheon
growing	1	Tyrell
strong	1	Tyrell

# Dictionary

**sorted so  
lookup is easy**

<b>coming</b>	<b>1</b>
<b>fury</b>	<b>1</b>
<b>growing</b>	<b>1</b>
<b>is</b>	<b>2</b>
<b>ours</b>	<b>1</b>
<b>strong</b>	<b>1</b>
<b>the</b>	<b>1</b>
<b>winter</b>	<b>1</b>

<b>Stark</b>
<b>Baratheon</b>
<b>Tyrell</b>
<b>Stark, Baratheon</b>
<b>Baratheon</b>
<b>Tyrell</b>
<b>Baratheon</b>
<b>Stark</b>

# Postings

<b>coming</b>	<b>1</b>		<b>Stark</b>
<b>fury</b>	<b>1</b>		<b>Baratheon</b>
<b>growing</b>	<b>1</b>		<b>Tyrell</b>
<b>is</b>	<b>2</b>		<b>Stark, Baratheon</b>
<b>ours</b>	<b>1</b>		<b>Baratheon</b>
<b>strong</b>	<b>1</b>		<b>Tyrell</b>
<b>the</b>	<b>1</b>		<b>Baratheon</b>
<b>winter</b>	<b>1</b>		<b>Stark</b>

# Search

coming	1	Stark
fury	1	Baratheon
growing	1	Tyrell
is	2	Stark, Baratheon
ours	1	Baratheon
strong	1	Tyrell
the	1	Baratheon
winter	1	Stark

winter

# Search

coming	1	Stark
fury	1	Baratheon
growing	1	Tyrell
is	2	Stark, Baratheon
ours	1	Baratheon
strong	1	Tyrell
the	1	Baratheon
winter	1	Stark

fury

# Search

coming	1	Stark
fury	1	Baratheon
growing	1	Tyrell
is	2	Stark, Baratheon
ours	1	Baratheon
strong	1	Tyrell
the	1	Baratheon
winter	1	Stark

is

# Search

<b>coming</b>	<b>1</b>	<b>Stark</b>
<b>fury</b>	<b>1</b>	<b>Baratheon</b>
<b>growing</b>	<b>1</b>	<b>Tyrell</b>
<b>is</b>	<b>2</b>	<b>Stark, Baratheon</b>
<b>ours</b>	<b>1</b>	<b>Baratheon</b>
<b>strong</b>	<b>1</b>	<b>Tyrell</b>
<b>the</b>	<b>1</b>	<b>Baratheon</b>
<b>winter</b>	<b>1</b>	<b>Stark</b>

**coming OR strong**



# Search

<b>coming</b>	<b>1</b>	<b>Stark</b>
<b>fury</b>	<b>1</b>	<b>Baratheon</b>
<b>growing</b>	<b>1</b>	<b>Tyrell</b>
<b>is</b>	<b>2</b>	<b>Stark, Baratheon</b>
<b>ours</b>	<b>1</b>	<b>Baratheon</b>
<b>strong</b>	<b>1</b>	<b>Tyrell</b>
<b>the</b>	<b>1</b>	<b>Baratheon</b>
<b>winter</b>	<b>1</b>	<b>Stark</b>

**fury AND growing**

# Searches Using Inverted Indices

Find all words ending with “ong”

strong



gnorts

Search for all words starting with “gno”

# Searches Using Inverted Indices

**Split words into n-grams for  
substring search**

yours



yo, you, our,  
ours, urs

**Match substrings with n-grams**

# Searches Using Inverted Indices

**Geo-hashes for geographical search**

**Algorithms such as Metaphone for  
phonetic matching**


**“Did you mean?” searches use a  
Levenshtein automaton**

An inverted index is at the  
heart of a search engine

# Implementing Search


---

# Apache Lucene



**The indexing and search library for a high performance, full-text search engine**


# Apache Lucene



**Open source, free to use**  
**written in Java, ported to other languages**




# Apache Lucene



Just like Hadoop in the distributed computing world, Lucene is the **nucleus** of several technologies built around it


# Apache Lucene



Solr

A **search server** with: distributed indexing,  
load balancing, replication, automated  
recover, centralized configuration


# Apache Lucene



**Nutch**

**Web crawling and index parsing**


# Apache Lucene



## CrateDB

**Open source, SQL distributed database**

# Elasticsearch



**Elasticsearch is a distributed search and analytics engine which runs on Lucene**


# Introducing Elasticsearch

---

# Elasticsearch



**An open source, search and analytics engine,  
written in Java built on Apache Lucene**



# Elasticsearch

**Distributed:** Scales to thousands of nodes

**High availability:** Multiple copies of data

**RESTful API:** CRUD, monitoring and other operation via simple JSON-based HTTP calls

**Powerful Query DSL:** Express complex queries simply


**Schemaless:** Index data without an explicit schema



# Elasticsearch



**Product catalog**  
**Inventory**  
**Autocomplete**



**Video clips**  
**Categories**  
**Tags**



**Courses**  
**Authors**  
**Topics**

# Elasticsearch



**Mining log data  
for insights**



**Price alerting  
platform**



**Business analytics  
and intelligence**

# Working with Elasticsearch

---

# Working with Elasticsearch



**As a service in the  
cloud**




**On your local  
machine**

# Working with Elasticsearch



**As a service in the  
cloud**



**On your local  
machine**

# Elasticsearch

<https://www.elastic.co/cloud/as-a-service>

**This however is subscription based, you can try it out free for 14 days.**

**This slide to be removed**

# Working with Elasticsearch



As a service in the  
cloud



**On your local  
machine**

Demo

**Download and install Elasticsearch on  
your local machine**



# Basic Concepts of Elasticsearch

---

# Near Realtime Search



**Very low latency, ~1 second from the time a document is indexed until it becomes searchable**

**Single server**  
**Performs indexing**  
**Allows search**  
**Has a unique id  
and name**

Node




# Cluster

Collection of nodes

Holds the **entire indexed data**

Has a **unique name**

Nodes **join** a cluster  
using the **cluster name**



# Document



A whole bunch of documents that need to  
be **indexed** so they can be **searched**

# Document



## catalog, reviews

# Document



**titles, description,  
comments**

Types



**Documents are divided into  
categories or *types***



Index



**All of these types of documents make up an **index****


**Collection** of similar documents

Identified by **name**

**Any number** of indices in a cluster

**Different indices** for different logical groupings

Index




**Logical partitioning of documents**

**User defined grouping semantics**

**Documents with the same fields belong to one type**

Type



# Document

Basic unit of  
information to be  
indexed

Expressed in **JSON**

Resides within an  
index


Assigned to a **type**  
within an index




# Documents in an Index




# Documents in an Index



# Documents in an Index



Too **large** to fit in the  
hard disk of one node



Too **slow** to serve all search  
requests from one node




Shards



**Split the index across  
multiple nodes in the cluster**




# Shards




## Sharding an index

Shards




**Search in parallel on  
multiple nodes**

# Replicas




Replicas




**High availability in case a  
node fails**

Replicas



**Scale search volume/throughput  
by searching multiple replicas**

# Shards and Replicas



An index can be split into **multiple shards**

A shard can be replicated **zero or more times**

An index in Elasticsearch has 5 shards and 1 replica by default

Demo

**Monitor the health of your cluster using  
HTTP requests**

# Summary

**Learnt a little search engine history,  
ubiquitous nature of search**

**Understood the basics steps involved in  
indexing and searching documents**

**Learnt how the inverted index data  
structure works**

**Got a brief introduction to Elasticsearch  
and its building blocks**

**Set up and installed Elasticsearch on  
your local machine**